Deployment

02476 Machine Learning Operations

Nicki Skafte Detlefsen,

Postdoc

DTU Compute

Loosely based on https://www.youtube.com/watch?v=2awmrMRf0dA

Freeing the model

MLOps

- Model deployment is part of the operations in MLOps
- In a nutshell: make the model available to others

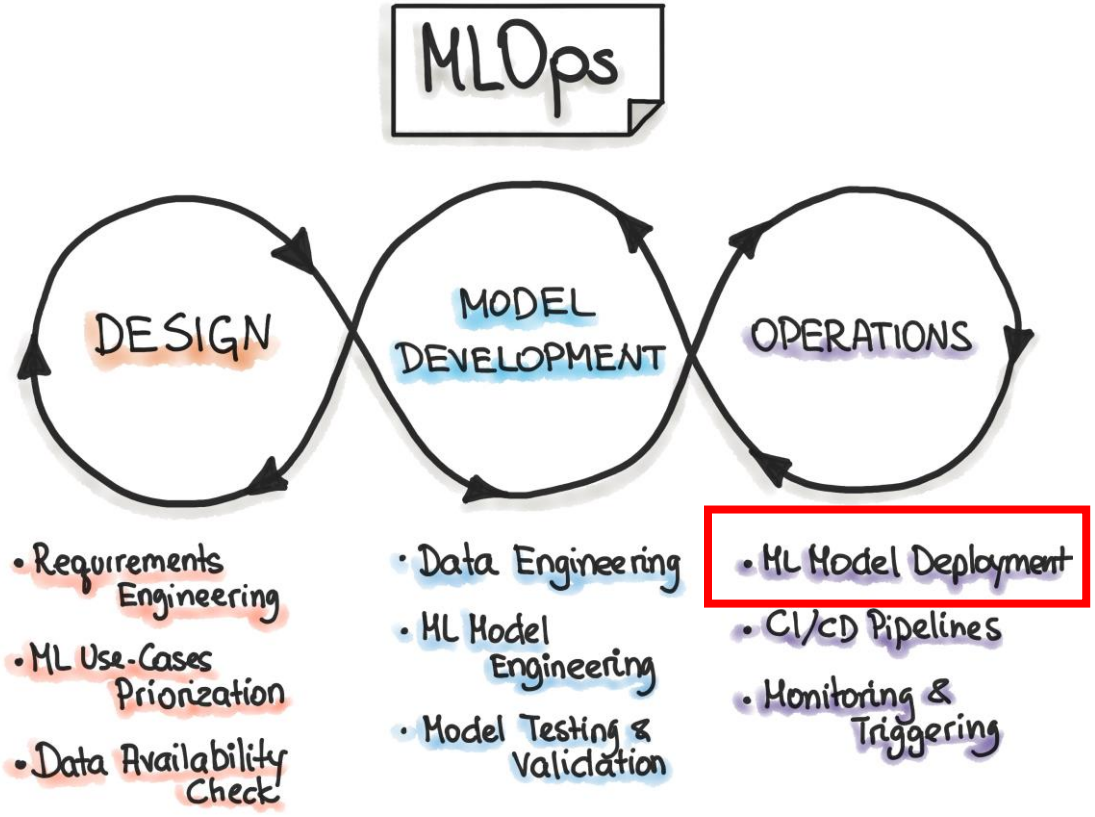

Remember this?

MLOps

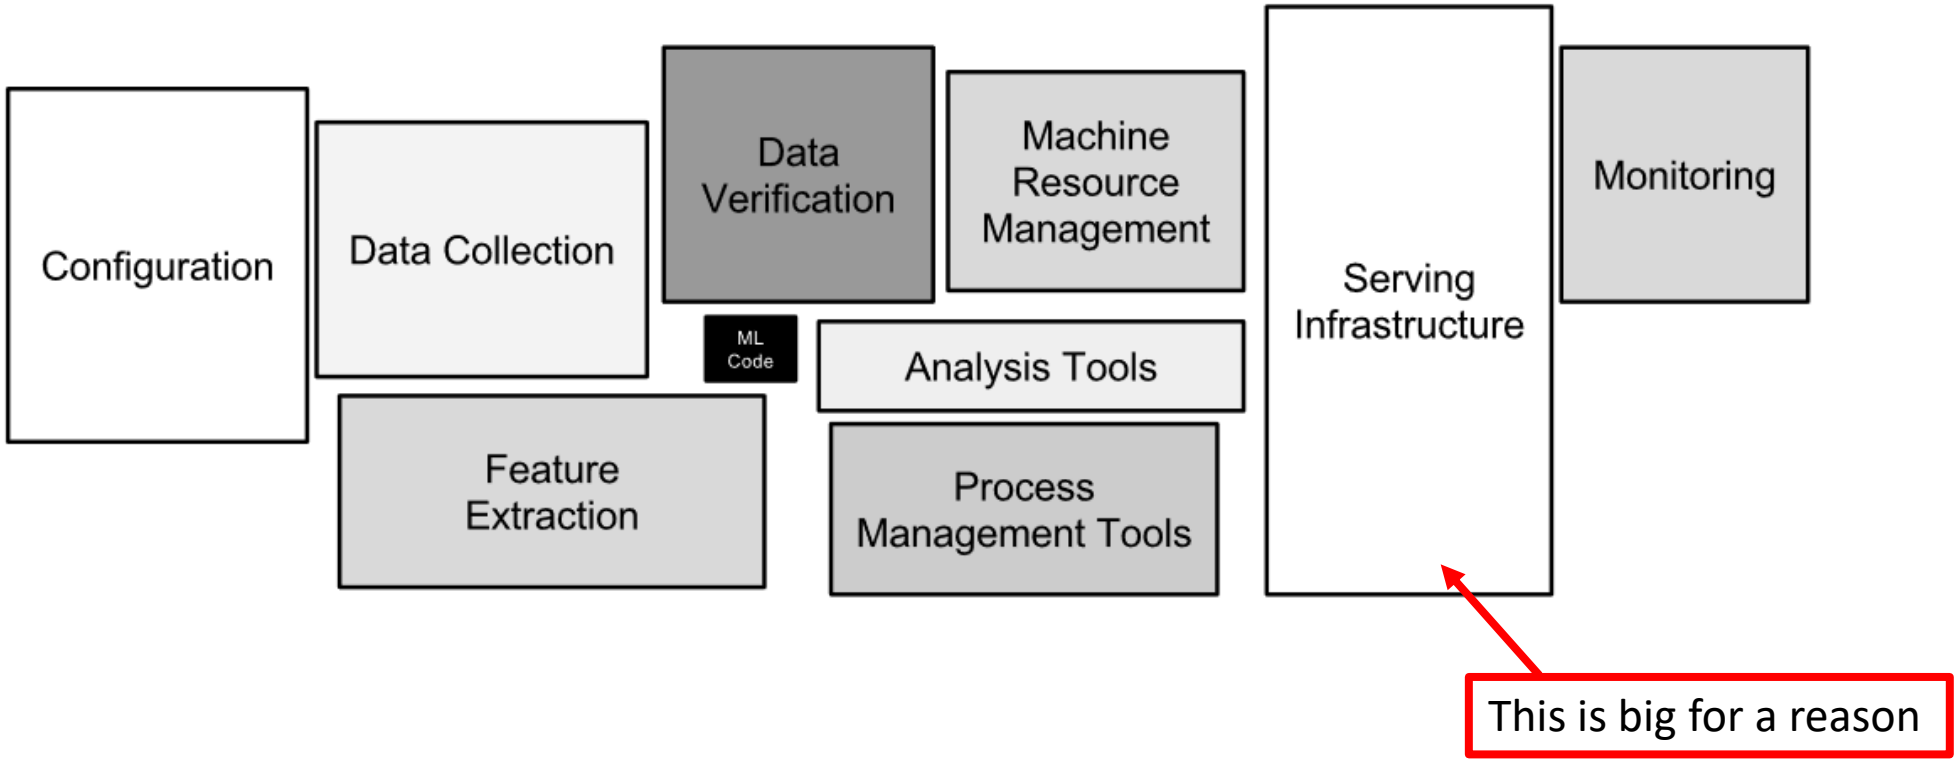

What do we want to deploy

In ML, *inferencing* refer to the use of a trained model to predict labels for new data on which the model has not been trained

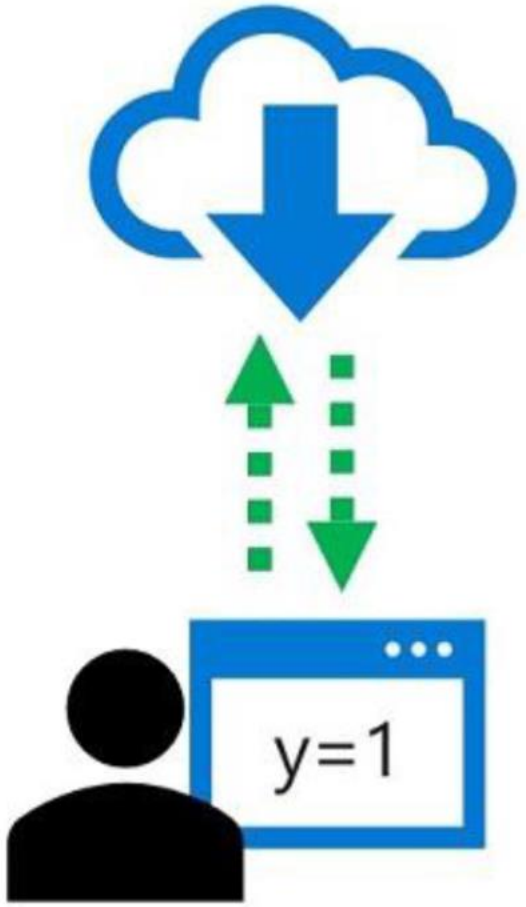

Many levels of deployment (within machine learning)

MLOps

1. Github reposatory + link to model weights
   - Easy to ”deploy”
   - Pain in the *** to use
2. Deploy on local computer/cluster
   - Fairly easy getting up and running, just requires people can access from outside
   - Can be fairly easy to use
   - Does not scale at all
3. Deploy to cloud service
   - Can be a pain to setup
   - Easy to use and scales to ∞ (and beyond!)

Production requirements

MLOps

1. Portability

   Models should be exportable to wide variety of enviroments, from C++ servers to mobile

2. Performance

   We want to optimize common patterns in neural networks to improve inference latency and throughput

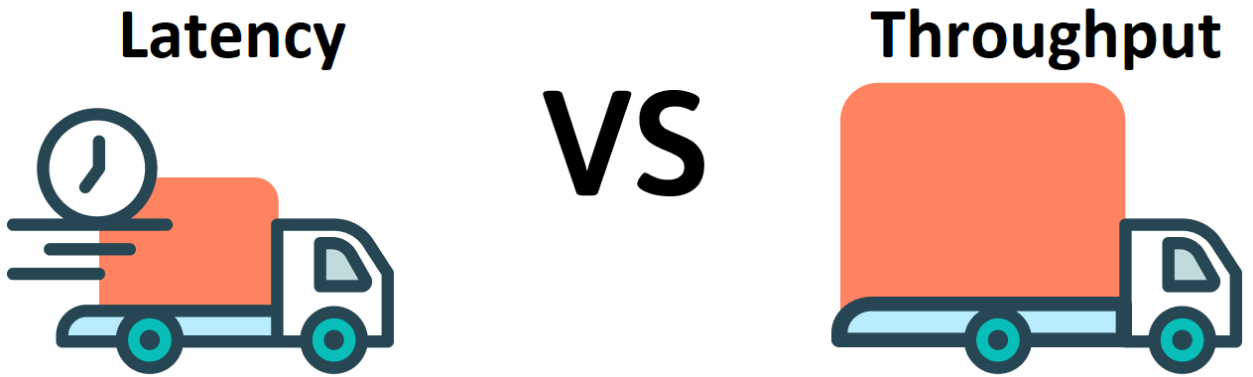

Choosing the right service

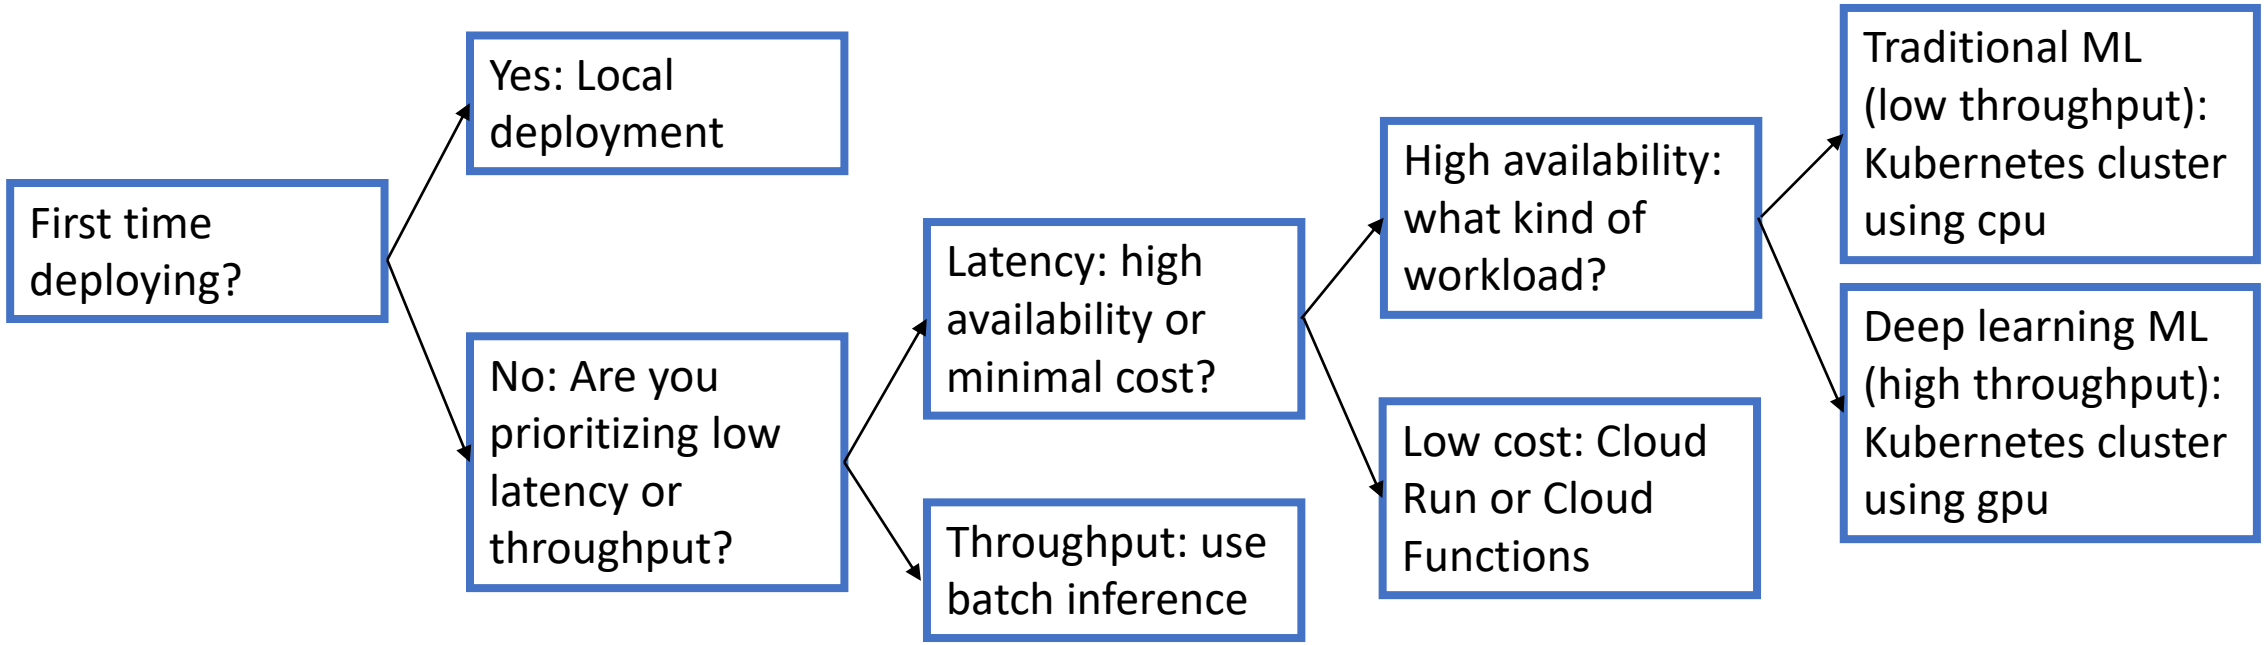

What are the challenges with Pytorch in production

- Pytorch is a dynamic framework (uses a dynamic graph)
  - This is not great in production as we need to know sizes etc. for compilation and optimization
- Why not use a static framework (Tensorflow 1.x, Caffe2 etc.)?
  - Do you really want to port all your work?
- What can we do to solve this?

Convert to script mode!

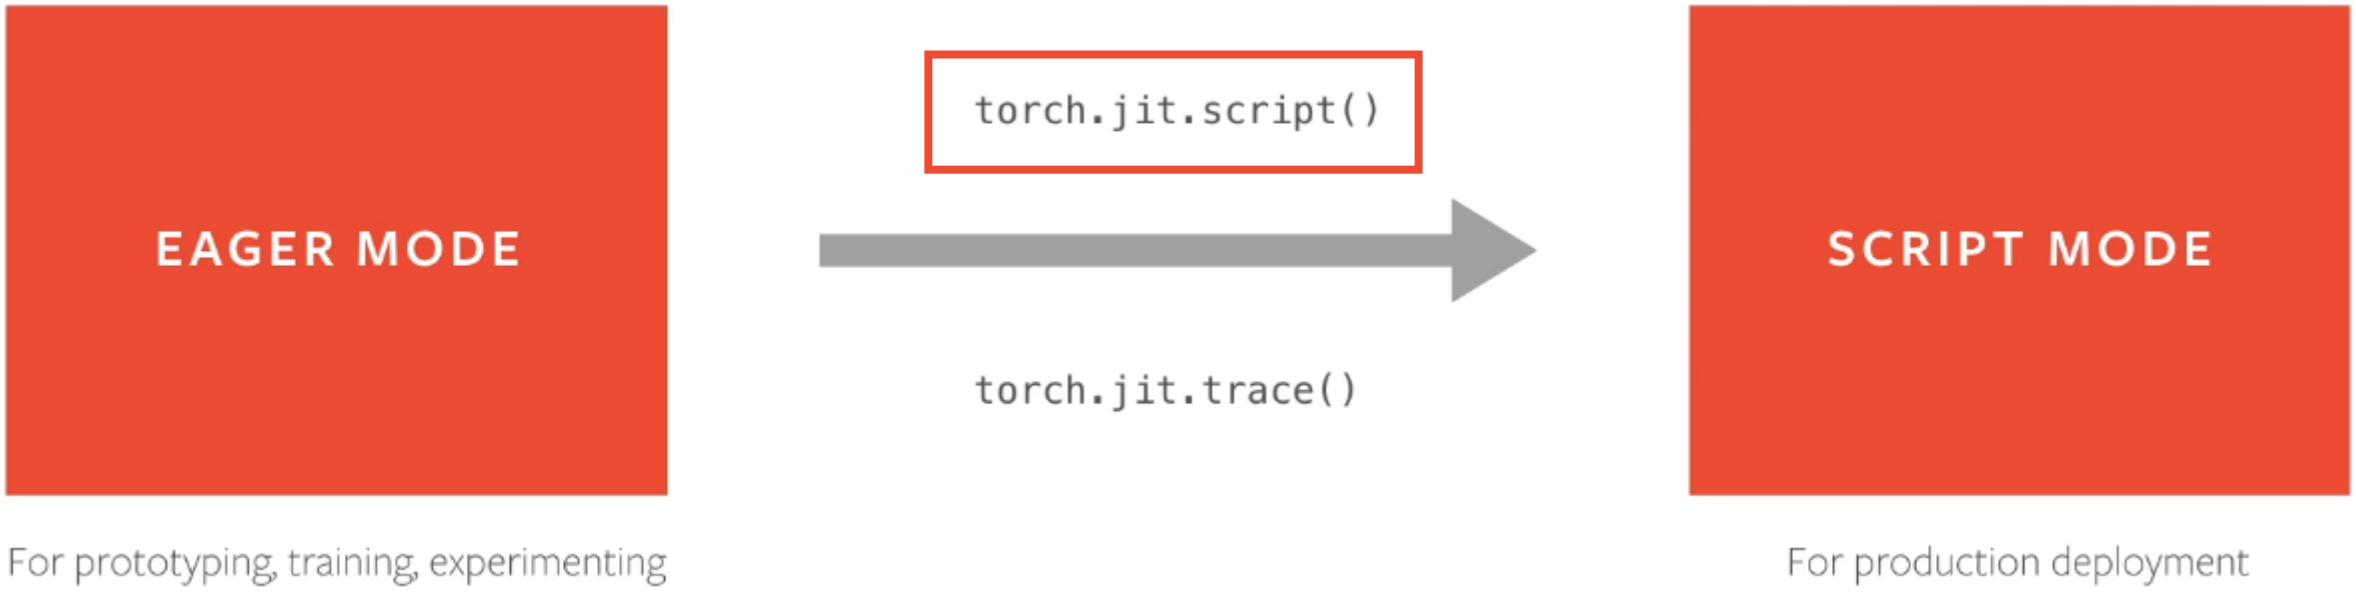

Serilization

- torch.jit.script **serialize the model, but what does it mean?**
- Serilization essentially encodes all modules methods, submodules, parameters, and attributes into a byte stream
- This makes the encoded model independent of python!
- This is basically just ”pickling” and “unpickling”.

Cloud functions

Simple one script files for deployment

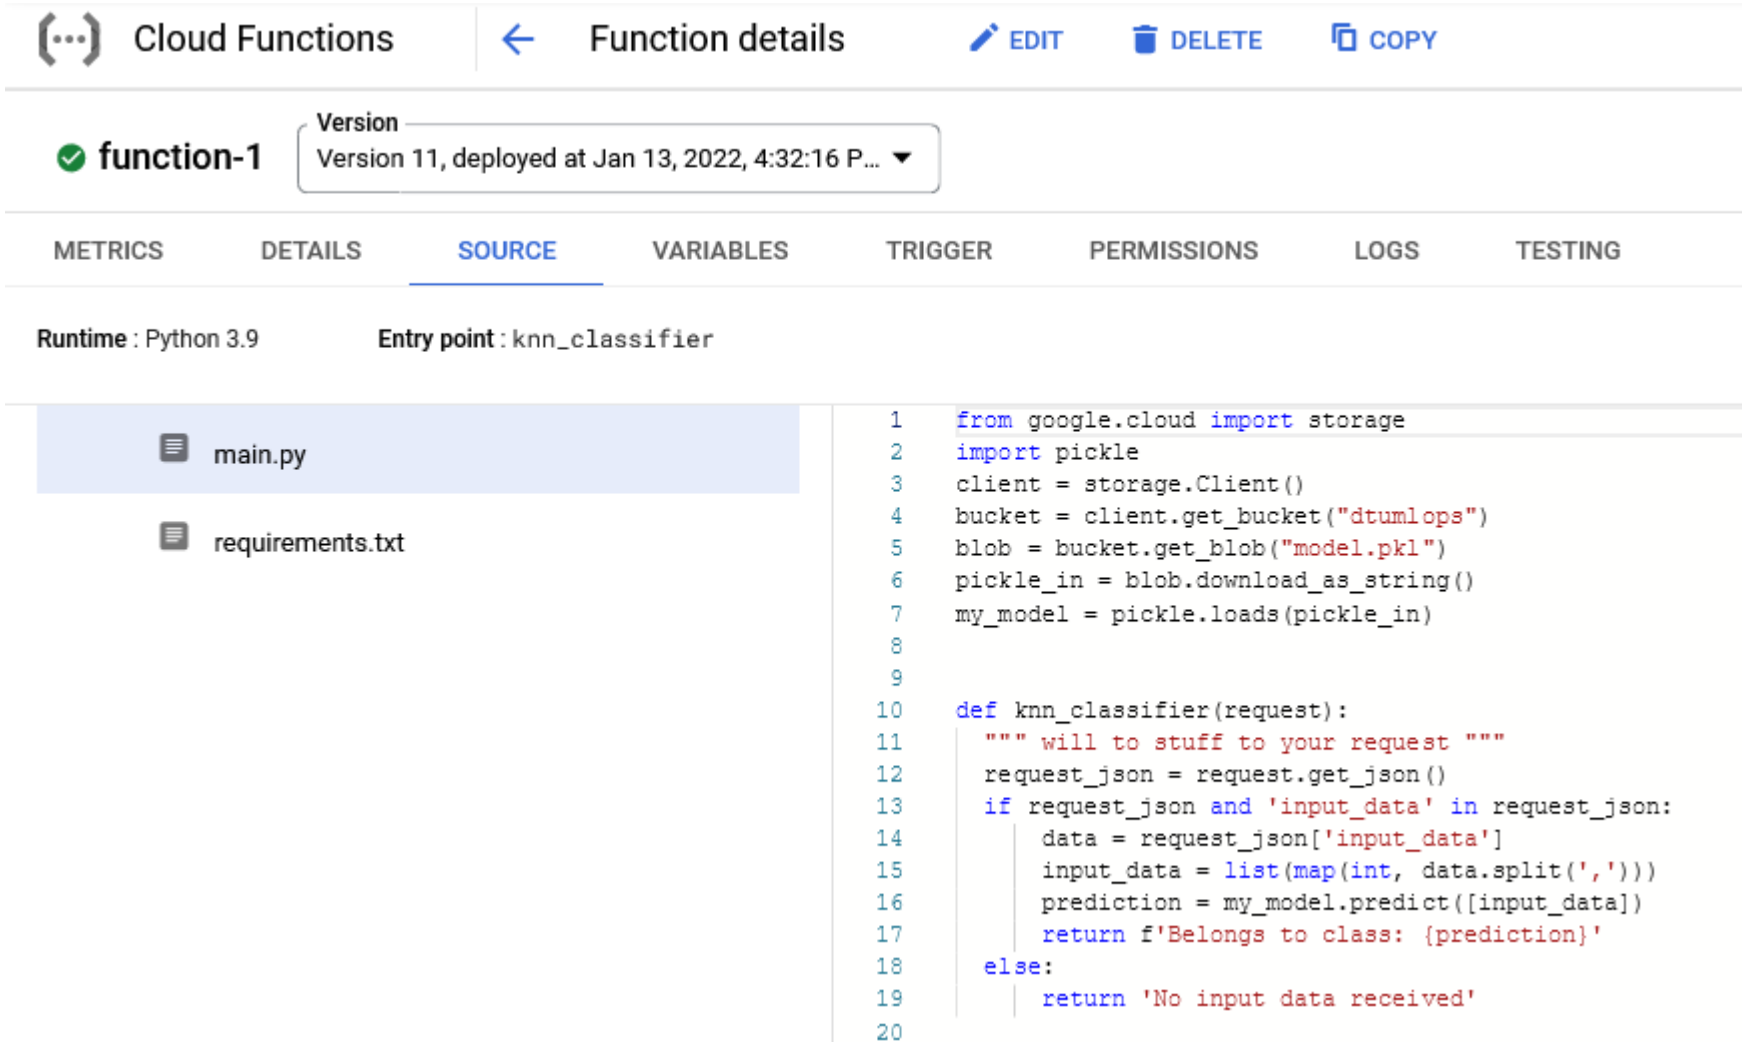

Meme of the day